

Συστήματα Αναμονής (Queuing Systems)

3η Εργαστηριακή Άσκηση

Λεούσης Σάββας

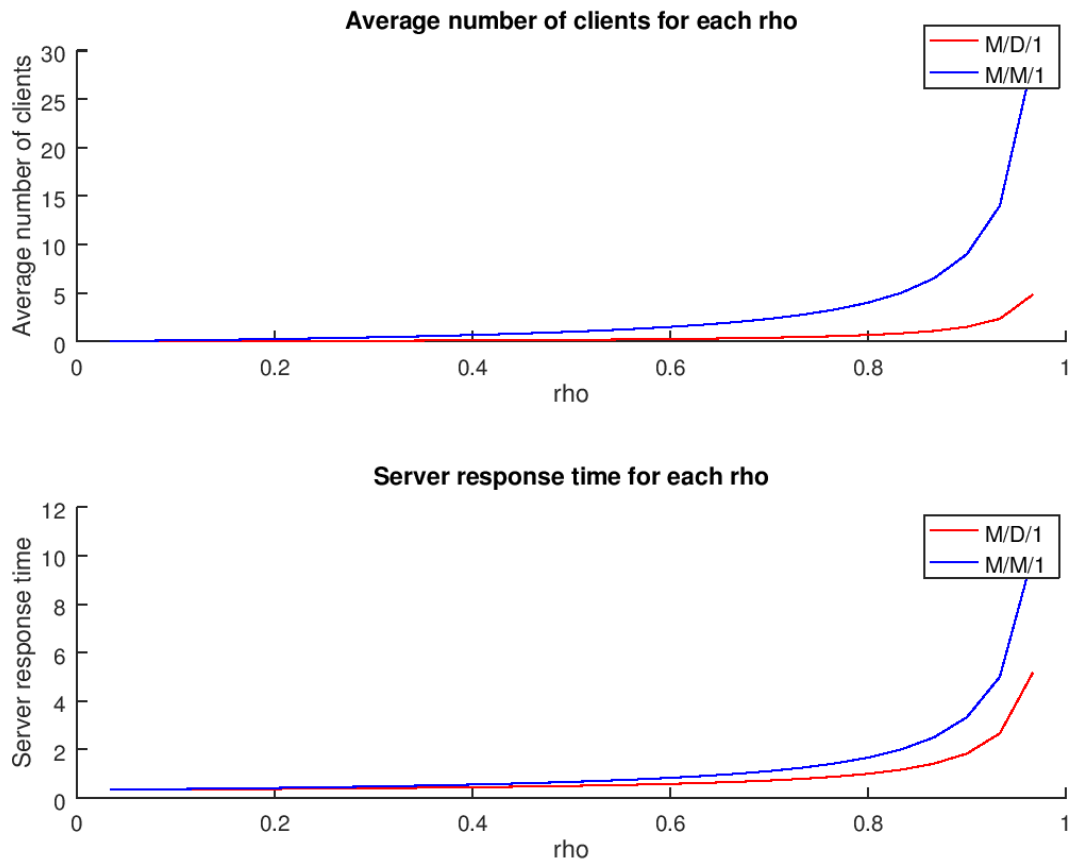
A.M.: 03114945

Σύγκριση συστημάτων M/M/1 και M/D/1

1. Ο μέσος χρόνος καθυστέρησης ενός πελάτη στην ουρά M/D/1, σύμφωνα με τον τύπο του Little, είναι $E(T) = \frac{E[n(t)]}{\lambda} = \frac{\rho + \frac{1}{2}(\frac{\rho^2}{1-\rho})}{\lambda} = \frac{\frac{2\rho - 2\rho^2 + \rho^2}{2-2\rho}}{\lambda} = \frac{2\rho - \rho^2}{(2-2\rho)\lambda} = \frac{\rho(2-\rho)}{2(1-\rho)\lambda}$ και ο μέσος χρόνος αναμονής είναι $E(W) = E(T) - \frac{1}{\mu} = \frac{\rho(2-\rho)}{2(1-\rho)\lambda} - \frac{1}{\mu} = \frac{\rho}{2(1-\rho)\mu} = \frac{\rho}{2(\mu-\lambda)}$, ενώ η απαραίτητη συνθήκη έτσι ώστε η ουρά M/D/1 να είναι εργοδική είναι να ισχύει ότι $\rho < 1 \Leftrightarrow \lambda < \mu$.
2. Η ζητούμενη συνάρτηση `qsmdd1.m` είναι για ουρές M/D/1 είναι η παρακάτω:

```
function [U R Q X] = qsmdd1( lambda, mu )
    if ( nargin != 2 )
        print_usage();
    endif
    ( isvector(lambda) && isvector(mu) ) || ...
        error( "lambda and mu must be vectors" );
    [ err lambda mu ] = common_size( lambda, mu );
    if ( err )
        error( "parameters are of incompatible size" );
    endif
    lambda = lambda(:)';
    mu = mu(:)';
    all( lambda >= 0 ) || ...
        error( "lambda must be >= 0" );
    all( mu > lambda ) || ...
        error( "The system is not ergodic" );
    U = rho = lambda ./ mu; % Server utilization
    R = (2-rho)./( (1-rho).*mu*2); % Server response time
    Q = rho./(2*(mu-lambda)); % Average number of requests
    in the system
    X = lambda; % Server throughput
endfunction
```

3. Χρησιμοποιώντας τις συναρτήσεις `qsmm1` και `qsmd1`, προκύπτουν οι παρακάτω γραφικές παραστάσεις:



Παρατηρώντας τις παραπάνω γραφικές παραστάσεις, βλέπουμε ότι από τις δύο ουρές τόσο ο μέσος αριθμός πελατών στο σύστημα, όσο και ο μέσος χρόνος καθυστέρησης είναι σαφώς μικρότεροι για την ουρά M/D/1. Επομένως, η ουρά M/D/1 είναι το καλύτερο σύστημα από τα δύο.

Προσομοίωση συστήματος M/M/1/10

Η ζητούμενος κώδικας προσομοίωσης ενός συστήματος M/M/1/10 είναι ο παρακάτω:

```
lambda = [1,5,10];
mu = 5;

for lambda = [1,5,10]
    total_arrivals = 0; % to measure the total number of
    arrivals
    current_state = 0; % holds the current state of the system
    previous_mean_clients = 0; % will help in the convergence
    test
    index = 0;
    clear arrivals;
    clear P;
    display(lambda);
    threshold = lambda/(lambda + mu); % the threshold used to
```

```

calculate probabilities
    rand("seed",1);
    transitions = 0; % holds the transitions of the simulation
in transitions steps

    while transitions >= 0
        transitions = transitions + 1; % one more transitions
step

        if mod(transitions,1000) == 0 % check for convergence
every 1000 transitions steps
            index = index + 1;
            for i=1:1:length(arrivals)
                P(i) = arrivals(i)/total_arrivals; % calculate the
probability of every state in the system
            endfor
            P_blocking = P(length(arrivals));
            mean_clients = 0; % calculate the mean number of
clients in the system
            for i=1:1:length(arrivals)
                mean_clients = mean_clients + (i-1).*P(i);
            endfor

            to_plot(index) = mean_clients;

            if abs(mean_clients - previous_mean_clients) < 0.00001
|| transitions > 1000000 % convergence test
                break;
            endif

            previous_mean_clients = mean_clients;

        endif

        random_number = rand(1); % generate a random number
(Uniform distribution)
#     if (transitions<=30) % debugging
#         display("##### NEW TRANSITION #####");
#         display(transitions);
#         display(current_state);
#         if current_state == 0 || random_number < threshold
#             display("Next transition is an arrival.");
#         else
#             display("Next transition is a departure.");
#         endif
#         display(total_arrivals);
#     endif
        if current_state == 0 || random_number < threshold %
arrival
            total_arrivals = total_arrivals + 1;
            try % to catch the exception if variable arrivals(i) is

```

```

undefined. Required only for systems with finite capacity.
    arrivals(current_state + 1) = arrivals(current_state
+ 1) + 1; % increase the number of arrivals in the current
state
    catch
        arrivals(current_state + 1) = 1;
    end
    if current_state == 10
        continue;
    else
        current_state = current_state + 1;
    endif
else % departure
    if current_state != 0 % no departure from an empty
system
        current_state = current_state - 1;
    endif
endif
endwhile

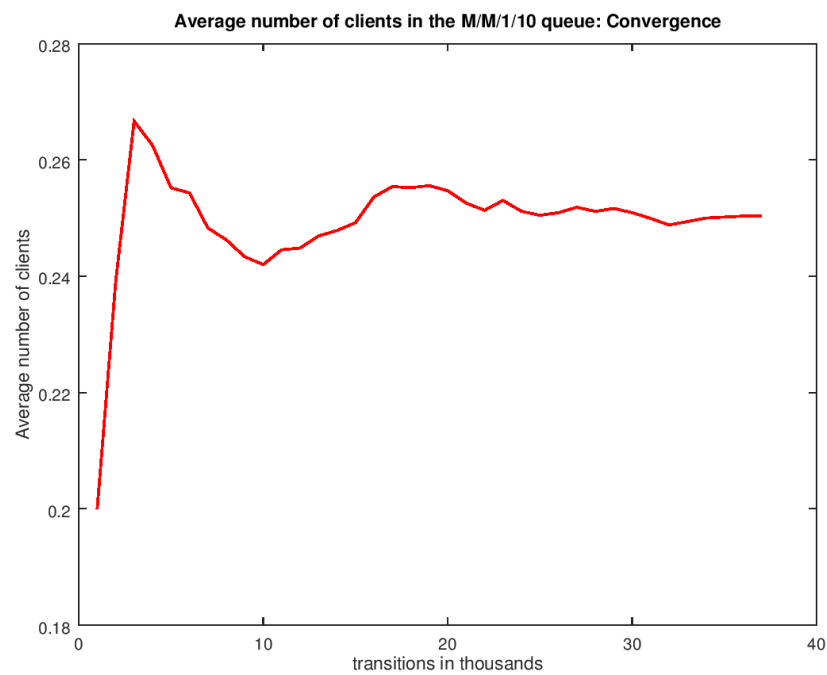

for i=1:1:length(arrivals)
    display(P(i));
endfor
display(P_blocking);
figure(fig_num++);
plot(to_plot,"r","linewidth",1.3);
title("Average number of clients in the M/M/1/10 queue:
Convergence");
xlabel("transitions in thousands");
ylabel("Average number of clients");

figure(fig_num++);
bar(P,'r',0.4);
title("Probabilities");
endfor

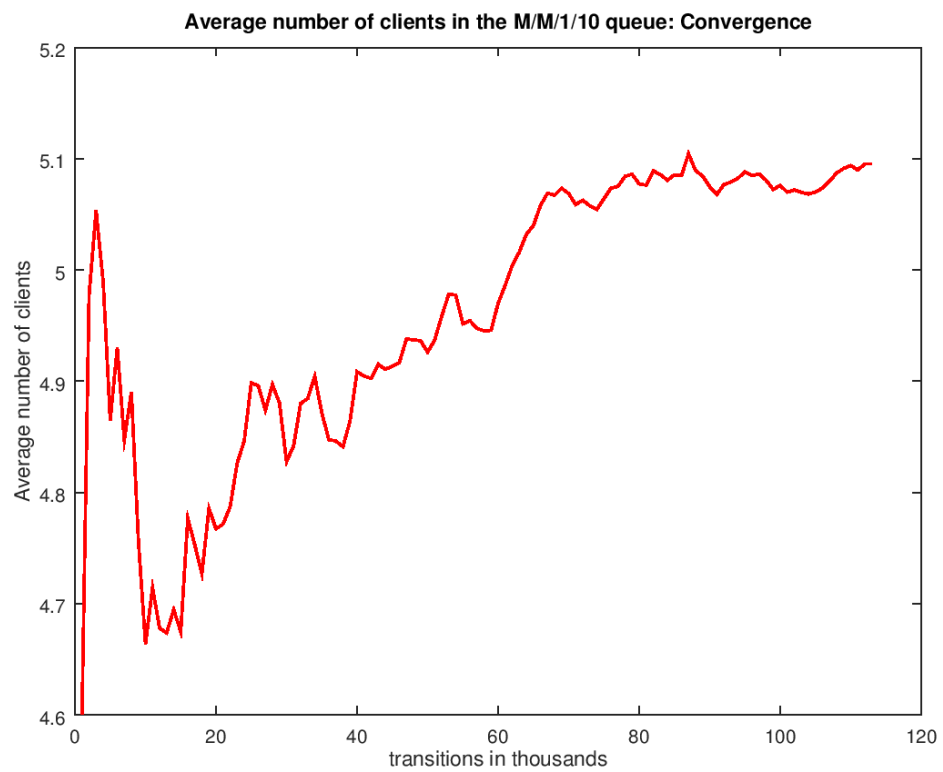
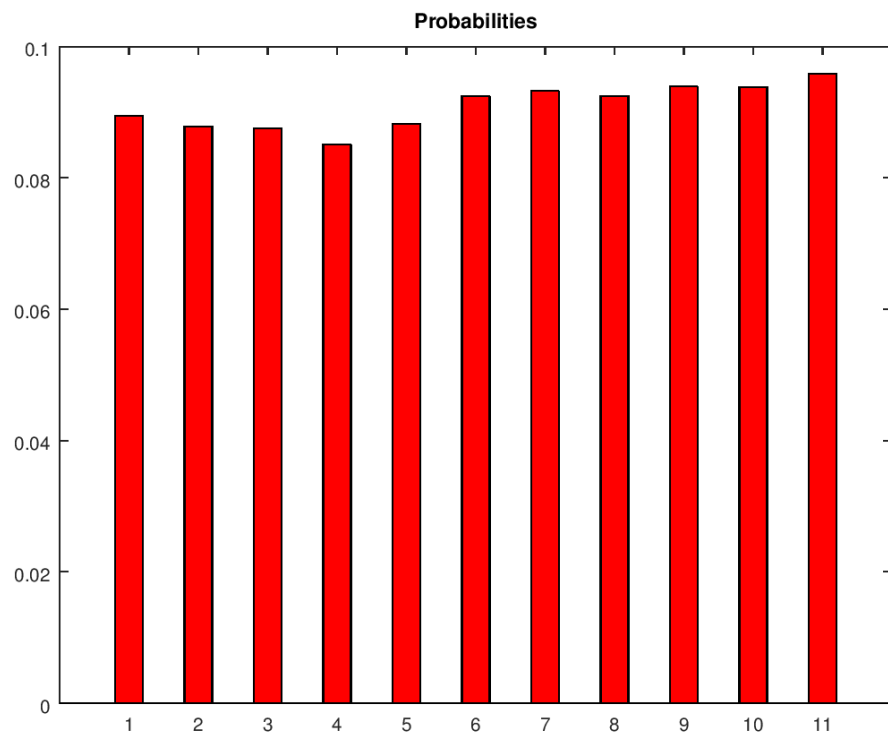
```

Για κάθε τιμή του λ προέκυψαν οι παρακάτω γραφικές παραστάσεις των εργοδικών πιθανοτήτων καθώς και της εξέλιξης του μέσου αριθμού πελατών στο σύστημα ανά 1000 μεταβάσεις:


- $\lambda=1$

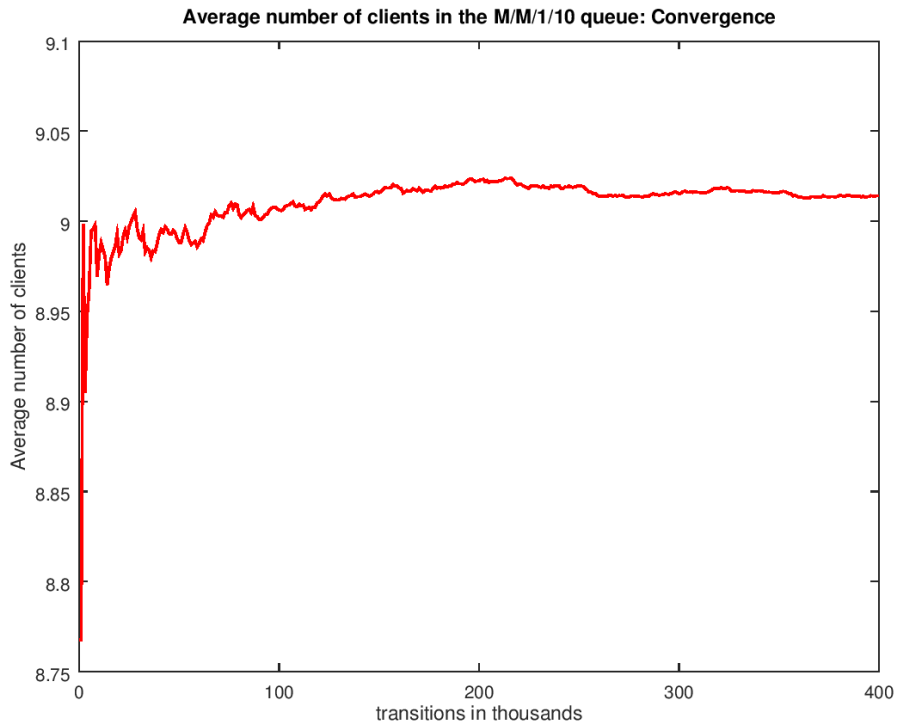


- $\lambda=5$



- $\lambda=10$





Σύμφωνα με τις παραπάνω γραφικές παραστάσεις, παρατηρούμε ότι όσο αυξάνεται το λ , τόσο αυξάνεται ο απαιτούμενος αριθμός μεταβάσεων έτσι ώστε το σύστημα να ικανοποιήσει το κριτήριο σύγκλισης. Προκειμένου να επιταχυνθεί η προσωμοίωση, θα μπορούσαμε να αγνοήσουμε τουλάχιστον 40.000 αρχικές μεταβάσεις, καθώς για τη σύγκλιση μας ενδιαφέρει η μόνιμη κατάσταση του συστήματος.

Προσομοίωση συστήματος M/M/1/5 με μεταβλητό μέσο ρυθμό εξυπηρέτησης

1. Οι εργοδικές πιθανότητες του συστήματος με τη βοήθεια του πακέτου queueing του Octave είναι οι εξής:

$P_0 =$

0.172822 0.252368 0.244132 0.177036 0.103107 0.050535

2. Τα ζητούμενα αποτελέσματα της προσωμοίωσης για την κάθε τιμή του κριτηρίου τερματισμού είναι τα παρακάτω:

a. 1%

```

criterion = 1
transitions = 2000
P =

```

0.174464 0.249513 0.237817 0.197856 0.091618 0.048733

```

P_blocking = 0.048733
mean_clients = 1.9288

```

b. 0.1%

```

criterion = 0.10000
transitions = 1000
P =

0.200000 0.262136 0.223301 0.161165 0.095146 0.058252

P_blocking = 0.058252
mean_clients = 1.8641

```

c. 0.01%

```

criterion = 0.010000
transitions = 1000
P =

0.200000 0.262136 0.223301 0.161165 0.095146 0.058252

P_blocking = 0.058252
mean_clients = 1.8641

```

d. 0.001%

```

criterion = 0.0010000
transitions = 1000
P =

0.200000 0.262136 0.223301 0.161165 0.095146 0.058252

P_blocking = 0.058252
mean_clients = 1.8641

```

e. 0.0001%

```

criterion = 1.00000e-004
transitions = 1000
P =

0.200000 0.262136 0.223301 0.161165 0.095146 0.058252

P_blocking = 0.058252
mean_clients = 1.8641

```

f. 0.00001%

```

criterion = 1.00000e-005
transitions = 1000
P =

0.200000 0.262136 0.223301 0.161165 0.095146 0.058252

P_blocking = 0.058252
mean_clients = 1.8641

```

g. 0.000001%

```

criterion = 1.00000e-006
transitions = 1000
P =

0.200000 0.262136 0.223301 0.161165 0.095146 0.058252

P_blocking = 0.058252
mean_clients = 1.8641

```

h. 0.0000001%

```

criterion = 1.0000e-007
transitions = 1000
P =

    0.200000    0.262136    0.223301    0.161165    0.095146    0.058252

P_blocking = 0.058252
mean_clients = 1.8641

```

Παράρτημα (κώδικας Lab3.m)

```

clc;
clear all;
close all;
fig_num = 1;
##### M/M/1 AND M/D/1 SYSTEMS COMPARISON
#####

# 3

lambda=0.1:0.1:2.9;
mu=[3];
[UD RD QD XD] = qsmd1(lambda,mu);
[UM RM QM XM p0] = qsmml(lambda,mu);
figure(fig_num++);
subplot(2,1,1);
hold on;
plot(lambda./mu,QD,"r");
plot(lambda./mu,QM,"b");
hold off;
title("Average number of clients for each rho");
xlabel("rho");
ylabel("Average number of clients");
legend("M/D/1","M/M/1");

subplot(2,1,2);
hold on;
plot(lambda./mu,RD,"r");
plot(lambda./mu,RM,"b");
hold off;
title("Server response time for each rho");
xlabel("rho");
ylabel("Server response time");
legend("M/D/1","M/M/1");

##### M/M/1/10 SYSTEM SIMULATION #####

lambda = [1,5,10];
mu = 5;

for lambda = [1,5,10]
    total_arrivals = 0; % to measure the total number of
    arrivals

```

```

    current_state = 0; % holds the current state of the
system
    previous_mean_clients = 0; % will help in the
convergence test
    index = 0;
    clear arrivals;
    clear P;
    display(lambda);
    threshold = lambda/(lambda + mu); % the threshold used
to calculate probabilities
    rand("seed",1);
    transitions = 0; % holds the transitions of the
simulation in transitions steps

    while transitions >= 0
        transitions = transitions + 1; % one more
transitions step

        if mod(transitions,1000) == 0 % check for
convergence every 1000 transitions steps
            index = index + 1;
            for i=1:length(arrivals)
                P(i) = arrivals(i)/total_arrivals; % calculate
the probability of every state in the system
            endfor
            P_blocking = P(length(arrivals));
            mean_clients = 0; % calculate the mean number of
clients in the system
            for i=1:length(arrivals)
                mean_clients = mean_clients + (i-1).*P(i);
            endfor

            to_plot(index) = mean_clients;

            if abs(mean_clients - previous_mean_clients) <
0.00001 || transitions > 1000000 % convergence test
                break;
            endif

            previous_mean_clients = mean_clients;

        endif

        random_number = rand(1); % generate a random number
(Uniform distribution)
        # if (transitions<=30) % debugging
        #     display("##### NEW TRANSITION #####");
        #     display(transitions);
        #     display(current_state);
        #     if current_state == 0 || random_number <
threshold
        #         display("Next transition is an arrival.");

```

```

#         else
#             display("Next transition is a departure.");
#         endif
#         display(total_arrivals);
#     endif
    if current_state == 0 || random_number < threshold %
arrival
        total_arrivals = total_arrivals + 1;
        try % to catch the exception if variable
arrivals(i) is undefined. Required only for systems with
finite capacity.
            arrivals(current_state + 1) =
arrivals(current_state + 1) + 1; % increase the number
of arrivals in the current state
        catch
            arrivals(current_state + 1) = 1;
        end
        if current_state == 10
            continue;
        else
            current_state = current_state + 1;
        endif
    else % departure
        if current_state != 0 % no departure from an empty
system
            current_state = current_state - 1;
        endif
    endif
endwhile

for i=1:1:length(arrivals)
    display(P(i));
endfor
display(P_blocking);
figure(fig_num++);
plot(to_plot,"r","linewidth",1.3);
title("Average number of clients in the M/M/1/10
queue: Convergence");
xlabel("transitions in thousands");
ylabel("Average number of clients");

figure(fig_num++);
bar(P,'r',0.4);
title("Probabilities");
endfor

##### M/M/1/5 SYSTEM SIMULATION WITH VARIABLE MU
#####

# 1

states = [0,1,2,3,4,5];

```

```

initial_state = [1,0,0,0,0,0];
lambda = 3;
mu = 1;
births_B = [lambda,lambda,lambda,lambda,lambda];
deaths_D = [mu*2,mu*3,mu*4,mu*5,mu*6];
transition_matrix = ctmcdb(births_B,deaths_D);
P = ctmc(transition_matrix);
for i=[1,2,3,4,5,6]
    index = 0;
    for T=0:0.01:50
        index = index + 1;
        P0 = ctmc(transition_matrix,T,initial_state);
        Prob0(index) = P0(i);
        if P0-P < 0.01
            break;
        endif
    endfor
endfor
display(P0);

# 2

lambda = 3;
mu = [1,2,3,4,5,6];

for criterion =
[1,0.1,0.01,0.001,0.0001,0.00001,0.000001,0.0000001]
    total_arrivals = 0; % to measure the total number of
arrivals
    current_state = 0; % holds the current state of the
system
    previous_mean_clients = 0; % will help in the
convergence test
    index = 0;
    clear arrivals;
    clear P;
    rand("seed",1);
    transitions = 0; % holds the transitions of the
simulation in transitions steps

    while transitions >= 0
        threshold = lambda/(lambda + current_state + 1);
        transitions = transitions + 1; % one more
transitions step
        if mod(transitions,1000) == 0 % check for
convergence every 1000 transitions steps
            index = index + 1;
            for i=1:length(arrivals)
                P(i) = arrivals(i)/total_arrivals; % calculate
the probability of every state in the system
            endfor
            P_blocking = P(length(arrivals));

```

```

        mean_clients = 0; % calculate the mean number of
clients in the system
        for i=1:1:length(arrivals)
            mean_clients = mean_clients + (i-1).*P(i);
        endfor

        to_plot(index) = mean_clients;
        if abs(mean_clients - previous_mean_clients) <
1/criterion || transitions > 1000000 % convergence test
            display(criterion);
            display(transitions);
            break;
        endif
        previous_mean_clients = mean_clients;
    endif
    random_number = rand(1); % generate a random number
(Uniform distribution)
    if current_state == 0 || random_number < threshold %
arrival
        total_arrivals = total_arrivals + 1;
        try % to catch the exception if variable
arrivals(i) is undefined. Required only for systems with
finite capacity.
            arrivals(current_state + 1) =
arrivals(current_state + 1) + 1; % increase the number
of arrivals in the current state
        catch
            arrivals(current_state + 1) = 1;
        end
        if current_state == 5
            continue;
        else
            current_state = current_state + 1;
        endif
    else % departure
        if current_state != 0 % no departure from an empty
system
            current_state = current_state - 1;
        endif
    endif
endwhile
display(P);
display(P_blocking);
display(mean_clients);
endfor

```